

# main

November 12, 2022

```
[766]: import pandas as pd
import nltk
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import BernoulliNB
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score
```

## 0.0.1 Step 1

```
[767]: df = pd.read_csv("federalist.csv")
df["author"] = df.author.astype("category")
print(df.head())
print()
df["author"].value_counts()
```

	author	text
0	HAMILTON	FEDERALIST. No. 1 General Introduction For the...
1	JAY	FEDERALIST No. 2 Concerning Dangers from Forei...
2	JAY	FEDERALIST No. 3 The Same Subject Continued (C...
3	JAY	FEDERALIST No. 4 The Same Subject Continued (C...
4	JAY	FEDERALIST No. 5 The Same Subject Continued (C...

```
[767]: HAMILTON          49
MADISON             15
HAMILTON OR MADISON 11
JAY                  5
HAMILTON AND MADISON 3
Name: author, dtype: int64
```

### 0.0.2 Step 2

```
[768]: X, y = df.text, df.author
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪train_size=0.8, random_state=1234)

print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

(66,) (66,)
(17,) (17,)
```

### 0.0.3 Step 3

```
[769]: stopwords = set(stopwords.words("english"))
vectorizer = TfidfVectorizer(stop_words=stopwords)
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

print(X_train.shape)
print(X_test.shape)

(66, 7876)
(17, 7876)
```

### 0.0.4 Step 4

```
[770]: nb = BernoulliNB()
nb.fit(X_train, y_train)
pred = nb.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, pred))
```

Accuracy Score: 0.5882352941176471

### 0.0.5 Step 5

```
[771]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪train_size=0.8, random_state=1234)
vectorizer = TfidfVectorizer(stop_words=stopwords, max_features=1000,
    ↪ngram_range=(1, 2))
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

nb.fit(X_train, y_train)
pred = nb.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, pred))
```

Accuracy Score: 0.9411764705882353

### 0.0.6 Step 6

```
[772]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳train_size=0.8, random_state=1234)
vectorizer = TfidfVectorizer(binary=True, stop_words=stopwords,
↳max_features=1000, ngram_range=(1, 2))
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

lr = LogisticRegression(solver="lbfgs", class_weight="balanced")
lr.fit(X_train, y_train)
pred = lr.predict(X_test)

print("Accuracy Score:", accuracy_score(y_test, pred))
```

Accuracy Score: 0.8823529411764706

### 0.0.7 Step 6 - model param change

I changed bigrams -> trigrams, introduced a min\_df, and increased the C value. But I did not see any improvements

```
[773]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳train_size=0.8, random_state=1234)
vectorizer = TfidfVectorizer(binary=True, stop_words=stopwords,
↳max_features=1000, ngram_range=(1, 3), min_df=5)
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

lr = LogisticRegression(solver="lbfgs", class_weight="balanced", C=2)
lr.fit(X_train, y_train)
pred = lr.predict(X_test)

print("Accuracy Score:", accuracy_score(y_test, pred))
```

Accuracy Score: 0.8823529411764706

### 0.0.8 Step 7

I tried a couple of layzer sizes but I couldn't find better results than a simple (15, 2) layer

```
[774]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳train_size=0.8, random_state=1234)
vectorizer = TfidfVectorizer(binary=True, stop_words=stopwords,
↳max_features=1000, ngram_range=(1, 2))
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)
```

```
nn = MLPClassifier(solver="lbfgs", alpha=1e-5, hidden_layer_sizes=(15,2),  
    ↪random_state=1)  
nn.fit(X_train, y_train)  
pred = nn.predict(X_test)  
  
print("Accuracy Score:", accuracy_score(y_test, pred))
```

Accuracy Score: 0.7058823529411765

[ ]: